

Does hyperbolic geometry help latent-space safety probes?

Llama-3-8B · harm monitoring (under adaptive jailbreak/obfuscation) and jailbreak-attack detection · negative result + a controlled twist

The question

- LLMs turn each prompt into ~4096 numbers (an **activation**). A **probe** reads them to flag harmful / jailbroken behavior.
 - Hyperbolic (**curved**) space is famously good at **hierarchies** — and harm has a hierarchy.
 - **Q: does measuring distance in curved space beat ordinary flat space for these latent-space safety probes?**
-

Two settings (so the terms are precise)

A. Harm monitor under adaptive attack (*Bailey "Obfuscated Activations"*)

- probe trained on **harmful vs benign**
- attacker **jailbreaks + hides the activation** to evade it

B. Attack-family detector (*JBSHield/RTV-style*)

- probe trained on **real jailbreaks** (GCG, PAIR, JBC) vs benign
- result: **C4 0.998 $\geq$ HPS 0.990 > HPS-Euclidean 0.968** (clean + open-set)

→ **Hyperbolic helps neither** (*A: clean+adaptive; B: clean+open-set — B-adaptive unrun*).

Main result — hyperbolic does NOT help (*Setting A: harm monitor under adaptive jailbreak/obfuscation*)



Clean + adaptive: rpf=C4, all bypassed under attack

- **Clean (left):** curved **rpf ties linear C4 at ~1.0** → curvature adds nothing.
 - Not just saturation: on the harder cache **rpf = rpfnoX0 = 0.998 exactly** (gap 0.000).
 - **Adaptive (right):** under attack, **every probe bypassed**; curved **tracks its flat twin**.
 - **Robust** across curvature and layers (*soft-prompt attack; GCG/HHD-adaptive unrun*).
-

Why — and it's NOT "linear is already perfect"

- "Linear is perfect → no room" only applies to **clean** detection.
- For the **adaptive** problem, curvature was a *genuinely plausible* fix — so **we tested it**:
 - under attack → bypassed like its flat twin;
 - in a *harder* open-set setting (held-out harm category) → curved beats flat-kNN (catch rate **0.77 vs 0.51**) **but still loses to the linear probe**.
- **Real reason = causal separability:** the attacker moves the activation across *whatever* boundary you draw, curved or straight. Geometry can't raise the attack cost.

- (One apparent "resistance" was **gradient masking** — a measurement artifact, not real robustness.)
-

What "rpf" is (the probes)



- Same hyperbolic pipeline; only the boxes in color differ.
 - **HPS** compresses 4096→64 then uses 12 hand-crafted features → the bottleneck (weak).
 - **rpf** keeps the full activation, no features → recovers linear-probe-level detection.
 - **rpfnox0** = rpf with the curved coordinate removed. **rpf** $\approx$ **rpfnox0** → **curvature is inert**.
-

Why clean works but adaptive fails (*Setting A — harm monitor, Bailey obfuscation game*)



PCA: harmful behaviors vs benign + obfuscated (jailbroken) activations

- Clean **harmful** (red) vs **benign** (blue) → **separate clusters**, a straight line splits them.
 - Orange = harmful **after the attack** → dragged **into the benign cluster**.
 - So the overlap is **adversarially created** → defeats every probe, any geometry.
 - (Note: harmful-vs-benign behaviors, not GCG/PAIR; orange = jailbroken-but-hidden.)
-

A different task — RETRIEVAL (*where hyperbolic finally helps; NOT jailbreak detection*)



Discrimination vs ranking

- Same hidden states, different job: not "is it a jailbreak?" but "**rank the most similar harmful prompts**" (retrieval).
 - **Left**: curvature *hurts* the decision (classification). **Right**: curvature *helps* retrieval (+0.13, and +0.16 with a deeper taxonomy).
 - **One knob trades discrimination for ranking** — curvature **HURTS** detection, **HELPS** retrieval — *faithful $\neq$ discriminative*.
-


Step 1 — turn prompts into a labeled tree



diagram 0

- 650 prompts → a **3-level path** each: category / leaf / sub-leaf .
 - 9 categories → 14 leaves → sub-leaves.
 - Tree distance = **how related**: same sub-leaf 0 · same leaf 2 · same category 4 · different 6.
 - *That distance is the ground truth the geometry should reproduce.*
-

Step 2 — train the space to SEPARATE categories



RANKING — curved (H) beats flat (E) at both depths (*immediate-parent retrieval, mAP, d=32*):

depth	curved H	flat E	gap
2-level	0.62	0.49	+0.13
3-level	0.52	0.37	+0.16

- **Robust claim: curved > flat at every depth.** Absolute mAP drops at 3-level (finer hierarchy = harder for both).
- "Grows with depth" is metric-dependent — the "parent" level itself shifts as the tree deepens (category→leaf), and at the very finest sub-leaf level the gap shrinks (+0.10→+0.06). So we call depth-scaling **suggestive, not a law** (MID slide). The solid result is just: **curved wins at every depth.**

DECIDING — tied at both depths; harm-vs-harmless won by a linear probe:

Job	metric	2-level (H/E)	3-level (H/E)
which category	top-1 / typed	0.79/0.74 · 0.72/0.73	0.70/0.66 · 0.54/0.53
harmful vs harmless (open-set kNN)	AUROC	linear 0.999 (curved 0.97, flat 0.95)*	—

- **Curved helps RANKING; ties on DECIDING; loses harm-vs-harmless to a plain linear probe.** That contrast IS the result.
- Depth-scaling: grows on harm same-target, mixed on MMLU → suggestive, not a law (MID slide). *open-set held-out-category, 2-level.

Grows with dimension — the evidence (*harm, mAP gap vs d*)

dimension d	2	4	8	16	32
curved – flat gap	−0.03	~0	+0.06	+0.10	+0.13

- At **low d (2–3) flat WINS**; curved only pulls ahead as **d grows**.
- **Textbook says the opposite** — hyperbolic is supposed to win at *low* d (cram a tree into 2–5 dims). Here it's reversed → **contrarian, novel**.
- *Likely why*: our reps are real high-dim LLM activations, not clean fixed-size graphs — but the **mechanism is still open** (honest caveat).

...and it's general, not a harm quirk



- The dissociation **replicates on MMLU subjects** — ground-truth labels, not harm.
- Curved beats flat at d=32: **+0.16 (category-level) / +0.07 (subject-level)** → **not harm-specific, not label-noise**.
- **Not surface-separable**: flat cosine is weak (0.41), unlike benign topics (0.86).

- This slide = **generality** (curved wins on a 2nd dataset), NOT depth-scaling (mixed on MMLU — MID slide).
-

Methodology contribution

- Standard curvature stats (δ -hyperbolicity, Ollivier-Ricci) are **unreliable at LLM dimension** — they call a sphere "hyperbolic."
 - We use a **calibrated** measure with a **dimension-matched random floor**.
 - Result: the **token/vocabulary space is strongly hyperbolic** (confirms HELM), but the **harm-decision direction is linear** → curvature exists, just not where the harm signal is.
-

● NOVEL — and why each is useful (*lit-search: 22 sources, per-claim verified*)

1. The discrimination-vs-ranking dissociation (flagship)

- one curvature knob, same reps: **hurts classification, helps retrieval**
- no prior work measures both axes on identical reps
- → *design rule*: **curve for ranking, stay flat for decisions**

2. The advantage GROWS with embedding dimension (contrarian)

- inverts the textbook "hyperbolic wins at low d" (De Sa'18, Bansal-Benton'21)
 - → *new fact about LLM geometry at scale*: curvature matters in **high-dim** reps
-

● MID — partly-new, and how to push each to novel

Grows with DEPTH — only confirmed on one dataset

- harm grows (+0.10→+0.16); MMLU flat (+0.065→+0.054)
- → fix: real **depth sweep (2→3→4)**, same target, ≥2 datasets

δ -hyperbolicity / Ollivier-Ricci unreliable at LLM dim

- new-in-this-form, but adjacent critiques exist
- → fix: ship a **standalone calibrated curvature protocol** + benchmark

Hyperbolic-retrieval-on-hidden-states (the setup)

- the combination is well-trodden (HIER, HyperbolicRAG, Raj)
- → fix: pitch the **dissociation + scaling**, not the setup

Raj (ICLR'26 workshop) — concurrent

- shares the premise, but classification-only, no dissociation/scaling
- → fix: cite + **distinguish on 4 axes**; our results are disjoint
- ⚠ **limitation it implies**: his hyperbolic win was *only* on **reasoning-distilled** models (late-layer compression); on standard Instruct (like our Llama-3) he found none → **our negative may not hold on a reasoning model (DeepSeek-R1)** — a next-step.

● *Known (cited, not claimed)*: Nickel-Kiela, HELM, Arditì'24, Bailey'24. Novelty claims are "to our knowledge" — they rest on a 22-source search.

Bottom line

1. **Hyperbolic does NOT help safety probes** — A (clean + adaptive) and B (clean + open-set). Linear wins. (B-adaptive unrun.)
 2. **But curvature DOES help hierarchical ranking** of harm content (+0.10–0.16 mAP). Not a detector.
 3. **NOVEL (to our knowledge): one knob trades discrimination for ranking, and the gain grows with dimension** — neither reported before.
-

Next steps / decision

Priority next experiment — attack the kNN detector

- only tested *clean* (0.97); run the adaptive attack on it
- if it stays above 0 where linear collapses → a **real robustness positive**
- (*clean score proves nothing* — linear is 0.999 *clean*, fully bypassed)

Then

- **publish now:** negative + dissociation + MMLU generality
 - **or invest:** 2nd model · data-size sweep (vs overfitting) · linear multi-class baseline
 - **done:** attack-family open-set (curved $\approx$ linear, above flat-kNN)
-

Q & A (anticipated)

Q: Is this jailbreak detection or harm detection? A: Setting A's probe is a **harm monitor** (trained on harmful vs benign) — the *jailbreak is the obfuscation* that evades it (the orange PCA points = the model jailbroken into complying while hidden). Setting B (real GCG/PAIR/JBC) is the actual **attack detector**. Hyperbolic helps neither.

Q: Would a probe trained on attacks raise the attack budget vs one trained on harm? A: Open — and **unrun**. It might help vs *known* attacks, but it overfits to attack signatures and causal separability lets the attacker move off them. We only tested it *clean* (~ 0.99), never under adaptive attack. Good next experiment.

Q: Why did rpf do better at $\lambda=0.5$? A: Noise — they cross back and forth (the flat twin leads at $\lambda=0.3$ and 0.7), all values near zero. Equal and both bypassed, not a win.

Q: Are the hidden states even hyperbolic? A: The token/vocabulary space is, strongly (matches HELM). But the harm-decision direction is linear — curvature exists, just not where the harm signal lives. Both true at once.

Q: Isn't "linear is perfect" the reason curvature can't help? A: Only for clean detection. For the adaptive problem curvature *could* have helped — that was the hypothesis — so we tested it; it doesn't, because the attacker crosses any boundary regardless of shape (causal separability).

Q: Is "hyperbolic on hidden states for retrieval" itself the novel part? A: **No**. That combination is well-trodden (HIER, HyperbolicRAG) + concurrent (Raj). Our novelty is on top of it: the **dissociation + grows-with-dimension**. We lead with those, not the setup.

Q: Good at harmful-vs-harmless but mediocre at harm categories — why? A: Opposite difficulty. Harmful-vs-harmless = **one big split** → a straight line nails it (0.999). Categories = **many fine splits among**

all-harmful items → subtle, less data each, noisy labels. The fine hierarchy is where geometry finally has structure to organize. *The gap IS the dissociation.*

Q: Is the retrieval positive a jailbreak result? A: No — it's organizing harmful content by category (ranking). Different task. Not a better detector.

Q: So what's it good for? A: Ranking/retrieval over hierarchies — triage, "find similar cases," taxonomy organization. Not detection.

Q: What exactly is rpf vs HPS? A: Same hyperbolic pipeline; HPS adds compression + 12 hand-crafted features (the bottleneck); rpf keeps the full activation; rpfnoX0 is rpf without the curved coordinate — and $rpf \approx rpfnoX0$, so curvature is inert.

Q: Is the positive big? A: Modest ($\sim +0.1$ – 0.2 mAP), ranking-only, grows with dimension (not the classic low-d story), one model. Real and significant, but modest.